

C++11 BLWS (Wednesday 1)

Smart Pointers

1. C++11: `std::shared_ptr`
 2. C++11: `std::weak_ptr`
 3. C++11: `std::unique_ptr`
 4. Deprecated `std::auto_ptr`
 5. Boost: Smart Pointer
 6. Boost: Scoped Pointer
 7. Boost: Scoped Array
 8. Boost: Intrusive Pointer
 9. Boost: Pointer Container
 10. Garbage Collection
-

Short breaks will be inserted as convenient.

C++11: `std::shared_ptr`

With `std::shared_ptr` C++11 introduced a **Smart Pointer** type that does reference counting on its pointee:

- When created by its default constructor it points to no object and creates an owner (reference) count set to zero.
- When created and initialized with a bare pointer it assumes it is the first and only one and creates an owner count set to 1.
- When initialized from another of its kind it points to the same object as the other (if any) and increments the – then shared – owner count.
- When assigned from another of its kind it first decrements the owner count and if that drops to zero, destroys the pointee; then it continues as if it were initialized from the pointer assigned to it.
- When it goes out of scope it also decrements the owner count and if that drops to zero, destroys the pointee.

Shared Pointee Construction

Per default an `std::shared_ptr` is initialized with no pointee.

Given

```
class MyClass { ... MyClass(bool, double, std::string); ... };
```

an `std::shared_ptr` can be initialized to point to a heap allocated object of type `MyClass` as follows:*

```
std::shared_ptr<MyClass> p1{new MyClass(true, 3.14, "hi!")};
```

Or:

```
auto p2 = std::make_shared<MyClass>(true, 3.14, "hi!");
```

The usual recommendation is to prefer the second way over the first as it can reserve space for `MyClass` together with the helper object (holding owner and observer count*) in a single heap allocation.

*: The observer count has not yet been introduced; its purpose will become clear when `std::weak_ptr` gets explained.

Shared Pointer Access

Given an `std::shared_ptr<SomeType> p` the most typical access to the pointee is via overloaded operator* or operator->, possibly after a testing whether there is an object:*

```
if (p) ... *p ...; // access whole object
      ... p->m ...; // access data member m
      ... p->f(); // call member function f
```

Furthermore `p.get()` returns the address of the pointee (or `nullptr`), so

- it can bridge between `std::shared_ptr` and legacy code that expects a native pointer,
- at least as long as the recipient is short-lived (compared to `p`) and
- **does not assume ownership.**

Shared Pointee Destruction

The default way to destruct the pointee (when the owner count drops to zero) is with `delete`.

If this is not appropriate a custom deleter can be specified at construction time:

```
std::shared_ptr<std::FILE> auto_close_fp{
    std::fopen("somefile", "r"),
    [](FILE *fp) { if (fp) std::fclose(fp); }
};
```

If the pointee is guaranteed to be valid or a custom deleter with a single pointer argument of pointee-type that is `nullptr`-safe* it could be specified directly:

```
if (auto fp = fopen("myfile", "w")) {
    std::shared_ptr<std::FILE> auto_close_fp{fp, std::fclose};
    ...
}
```

: Implementations of `fclose` differ in C, some are more robust and check the pointer argument for `(FILE)0`, but according to the C-Standard an invalid pointer causes undefined behavior.

C++11: `std::weak_ptr`

With `std::weak_ptr` C++ introduced a companion to `std::shared_ptr`, mainly used to break cyclic references, which would otherwise defeat one of the main motivations for using smart pointers as a light-weight, high-efficiency garbage collector.

- An `std::weak_ptr` acts as **observer** of a pointee **owned** by an `std::shared_ptr`.
- As such it shares and manages an observer count (similar to but different from the owner count).
- A non-zero observer count **will not keep the pointee alive** if the owner count drops to zero.
- Therefore an `std::weak_ptr` has no way to access the pointee directly via overloaded `operator*` or `operator->`.
- To gain access it has first to obtain an `std::shared_ptr` – which might fail but if successful will keep the pointee alive even if all other owners cease to exist.

Weak Pointee Usage

Per default an `std::weak_ptr` is initialized with no pointee.

Given `std::shared_ptr<MyClass> p` an `std::weak_ptr` can be initialized to the pointee referred by `p` (if any) with:

```
std::weak_ptr<MyClass> wp{p};
```

To get access to the pointee an `std::shared_ptr` must be obtained and tested:

```
if (auto sp = wp.lock()) {  
    // sp != nullptr  
    // now owns the object wp had observed  
    ... *sp ...;    // access whole object  
    ... sp.m ...;   // access data member m  
    ... sp->f();    // call member function f  
} // sp goes out of scope, if all other owners  
  // are gone pointee will get destroyed here
```

C++11: `std::unique_ptr`

An `std::unique_ptr` is – as the name suggests – the sole owner of its pointee:

- Therefore there can always be only one for each pointee, i.e.
 - it is **not possible** to copy-construct or copy-assign an `std::unique_ptr` from another one of its kind (yields a compile time error), but
 - it **is possible** to move-construct or move-assign an `std::unique_ptr` from another one of its kind.
- When an `std::unique_ptr` goes out of scope or is re-assigned it first destroys its pointee (if any).

The implementation of `std::unique_ptr` is very close to native pointers, i.e. same memory footprint, and also for most operations same performance, except those that need to care for destruction of a (previous) pointee.

Unique Pointee Construction

Per default an `std::unique_ptr` is initialized with no pointee.

Given any type `T` (a builtin type, a class from the standard library, or a user defined class) an `std::unique_ptr` can be initialized to point

- to a **single heap allocated object of this type***

```
std::unique_ptr<T> ptr{new T};
```

- or to an **array of N heap allocated objects**:

```
std::unique_ptr<T[]> arr{new T[N]};
```



The appropriate deleter is set depending on the template instantiation argument `T` or `T[]`.

*: Constructor arguments may be supplied as usual but there is no `std::make_unique` analogous to `std::make_shared` until C++14.

Unique Pointer Access

Given an `std::unique_ptr<SomeType> p` the most typical access to the pointee is via overloaded `operator*`, `operator->`, or `operator[]` if the pointee is an array, possibly after a testing whether there is an object:*

```
if (p) ... *p ...; // access whole object
... p->m ...; // access data member m
... p->f(); // call member function f
... p[i] ...; // access i-th whole object (0-origin)
```

Furthermore

- `p.get()` returns the address of the pointee (or `nullptr`), so it can bridge between `std::unique_ptr` and legacy code that expects a native pointer with shorter lifetime as `p`;
- `p.release()` is similar but relinquishes ownership of the pointee **which the recipient has to assume then.**

Moving Unique Pointers

As unique pointers can not be copied, the (deliberate) use as initial values for copy constructor arguments is not possible:

```
extern std::unique_ptr<MyClass> make_MyClass();  
...  
std::unique_ptr<MyClass> p1 = make_MyClass(); // OK (move c'tor)  
std::unique_ptr<MyClass> p2(p1); // ERROR (no copy c'tor exists)
```

Instead the existing pointer needs to be moved:

```
std::unique_ptr<MyClass> p3(std::move(p1)) // OK (p1 nullptr now!)
```

Same for assignment – move works, copy does not:

```
p1 = p3; // ERROR (no copy assignment exists)  
p1 = make_MyClass(); // OK (move assignment)  
p1 = std::move(p3); // OK (p3 nullptr now!)
```

Unique Pointee Destruction

Per default the pointee is destructed with `delete` or `delete[]` depending on the way an `std::unique_ptr` has been created. Wrong pairing will not be detected at compile time but cause undefined runtime behaviour.*

```
{ std::unique_ptr<T[]> ptr{new T};  
    ...  
} // destructor does delete[] on pointee address
```

Or:

```
{ std::unique_ptr<T> arr{new[N] T};  
    ...  
} // destructor does delete on pointee address
```



Pairing plain allocation with array deallocation or array allocation with plain deallocation has undefined behavior.

*: In the best case this will cause an immediate crash with a good error message. But a crash may also occur much later with a misleading error message (if any) and therefore may be hard to relate to its original cause, or there may be a memory leak, memory overwritten with bad values, whatever ...

Deprecated `std::auto_ptr`

The only **Smart Pointer** in C++98 was `std::auto_ptr`, which is deprecated since C++11.

- It had nearly the same behavior (and implementation) as `std::unique_ptr` has now, but C++98 had no means to forbid the copying versions of constructor and assignment while still allowing the move versions.
- Therefore `std::auto_ptr` had copy-constructor and -assignment which set their right hand side to `NULL`, i.e. the auto-pointer used for initialisation or from which the pointee was assigned lost its pointee.
- **That came as a surprise to some developers who expected a "more intelligent" behaviour from a "smart pointer".**

Or as Bjarne Stroustrup once put it: With C++11 `std::unique_ptr` became what `std::auto_ptr` in C++98 always should have been but couldn't, due to lacking proper language support.

Boost: Smart Pointer

The term *Smart Pointer* is also used to subsume pointer-like helper classes in Boost:

- `boost::shared_ptr`:
much like `std::shared_ptr` (the latter mostly emerged from the former);
- `boost::weak_ptr`:
much like `std::weak_ptr` (the latter mostly emerged from the former);
- `boost::scoped_ptr` and `boost::scoped_array`:
close to `std::unique_ptr` but in two variants to provide different destructors to do a plain `delete` or a `delete[]` on the pointee.
- `boost::intrusive_ptr`:
similar in purpose to `boost::shared_ptr` / `std::shared_ptr` but storing the reference count inside the pointee (which therefore must be accessible).

Boost: Scoped Pointer

A **Scoped Pointer** considers itself to be the sole owner of a

- **single object allocated on the heap**

and will finally destroy its pointee (if any) when going out of scope or a new pointee is assigned.

- There is no copy-constructor and -assignment, the only way to re-assign a `boost::scoped_ptr` is via `swap` (available globally and as member function);
- Final destruction will use `delete`, therefore expecting the pointee is a single object.

[!] If a `boost::scoped_ptr` is initialized with an address not returned from `new` or is pointing to an array of objects returned from array heap allocation, undefined behaviour will result at deletion time.

Boost: Scoped Array

A **Scoped Array** considers itself to be the sole owner of

- **an array of objects allocated on the heap**

and will finally destroy its pointee (if any) when it goes out of scope or a new pointee is assigned.

- There is no copy-constructor and -assignment, the only way to re-assign a `boost::scoped_array` is via `swap` (available globally and as member function);
- Final destruction will use `delete[]`, therefore expecting the pointee is an array of objects.



If a `boost::scoped_array` is initialized with an address not returned from `new` or pointing to a single object returned from plain heap allocation, undefined behaviour will result at deletion time.

Boost: Intrusive Pointer

An **Intrusive Pointer** is much like a reference counted `boost::shared_ptr` or `std::shared_ptr`.

- Instead of allocating reference counts separately it expects two global functions overladed for pointers to the pointee's type:
 - `intrusive_ptr_add_ref` – called when a new referrer for the pointee is added.
 - `intrusive_ptr_release` – called when an existing referrer of the pointee gets re-assigned or goes out of scope.
- Furthermore there is a class `boost::intrusive_ref_counter` from which the pointee's class may be derived.*

Boost recommends in case of doubt to prefer ordinary shared pointers and to avoid using intrusive pointers without good reason.

*: Of course given its source is written from scratch or at least available and can be modified.

Intrusive Pointer Example

To make MyClass usable with intrusive pointers it can be written as:*

```
class MyClass :  
    public boost::intrusive_ref_counter<  
        MyClass,  
        boost::thread_unsafe_counter  
    > { ... };
```

Or if it should be usable in a multi-threaded environment:

```
class MyClass  
    : public boost::intrusive_ref_counter<  
        MyClass,  
        boost::thread_safe_counter  
    > { ... };
```

Then there can be intrusive pointers of MyClass:

```
boost::intrusive_ptr<MyClass> p{new MyClass(...)};
```

*: The second template argument may be omitted as it defaults to boost::thread_unsafe_counter.

Boost: Pointer Container

A number of [Pointer Containers] has been made available by boost, paralleling the STL containers with a pointer version that

- omits the pointer syntax at instantiation,
- adds one level of dereferencing to each member access, and
- considers its elements as pointers owning the memory pointed to.

While the former two are more a matter of convenience (see example on next slide), the last one has a severe semantical implication:

- If the container goes out of scope it deletes all the pointees of its (still) contained elements.*

[!] Storing non-owning pointers or pointers that do not even point to heap-allocated memory in a pointer container will cause undefined behaviour when or after the container goes out of scope.

*: This effect can alternatively be achieved by storing `std::unique_ptr`-s in an ordinary container.

Pointer Container Example

Storing and later on processing a `boost::ptr_vector`:^{*}

```
boost::ptr_vector<MyClass> v;
...
// fill in some content (probably in a loop):
... v.push_back(new MyClass(...));

// process later (or maybe in a different thread):
while (!v.empty()) {
    ... v.back() ...; // access MyClass as a whole
    ... v.back().m ...; // access MyClass data member
    ... v.back().f(); // call MyClass member function
    v.pop_back();
}
```

In case the processing loop is not reached or left before the content is fully processed, **the pointer container destructor** will call delete for the pointers still contained, avoiding a memory leak.

^{*}: If a pointer container actually gets filled and processed concurrently as suggested by the comment in the example, mutexes or other synchronization techniques must be added as modifying operations are not thread-safe by themselves.

Pointer Container Substitute

Storing and later on processing a container of custodial pointers:*

```
std::vector<std::unique_ptr<MyClass>> v;
...
// fill in some content (probably in a loop):
... v.push_back(new MyClass(...));

// process later (or maybe in a different thread):
while (!v.empty()) {
    ... *(v.back()) ... // access MyClass as a whole
    ... v.back()->m ... // access MyClass data member
    ... v.back()->f() ... // call MyClass member function
    v.pop_back();
}
```

In case the processing loop is not reached or left before the content is fully processed, **the `std::unique_ptr` destructors** will call delete for their pointees, avoiding a memory leak.

*: If an STL container actually gets filled and processed concurrently as suggested by the comment in the example, mutexes or other synchronization techniques must be added as modifying operations are not thread-safe by themselves.

Garbage Collection

There is no garbage collection in C++ because of a specific difficulty:

- **An address once obtained from `new` may not be visible in any memory location capable of holding a heap address**, instead it
 1. may have been modified by address arithmetic ... which will of course be reverted before the `delete` takes place;
 2. may be temporarily stored in an integral type^{*} ... and will of course be restored to a pointer of appropriate type before the `delete` takes place.

Both are not a sign of bad programming style but have some valid uses in the C and C++ code base written in the last 35 years, so they cannot be easily ruled out by a new language standard.

^{*}: C/C++ even guarantees that when an integral type of sufficient size is used as temporary to store a pointer, after assigning the content back to the original pointer type the memory location pointed to will not have changed ... which by no means says that the bit patterns stays the same all the time!

C++11: Garbage Collection API

C++11 has defined an [Garbage Collection ABI](#) to enable *Interested Third Parties* to supply garbage collection as add-on library.*

Mainly the ABI allows to say (put colloquially):

- **The object at this address I name to you may appear not to be any longer in use.**
 - You may not find it in any memory location capable of holding a heap pointer. Nevertheless be assured: it *is* still in use, so *do not garbage collect it*, I'll take up responsibility and return the reserved space in due course when its *really* not in use any more.
- **In this memory area I name to you, you may find storage cells looking like pointers to heap memory, but they aren't.**
 - So, in case there is any memory pending to be freed and *its only use appears to be from inside this area*, feel free to go ahead and garbage collect that stuff.

*: It will surely be interesting to watch such efforts and whether any third-party garbage collector for C++ gets into wide-spread use. If so, then probably rather for new software, not for large amounts of legacy code (including libraries), and maybe only with additional support by compiler warnings.